

Individual Project Practical Checkpoint 1

Checkpoint report template

Replace each gray prompt with concise evidence from your current project. Keep the completed report within two pages and export it as a PDF. Submit the working technical artifact separately.

1 Project Direction


Field	Response
Student	Rickey Johnson
Working title	Locksmith: Comparing Prompt-Injection Defense Strategies Through Adversarial Gameplay
Approved question or objective	Does raising defense-prompt complexity across levels measurably lower extraction success?
Changes from feedback	No changes.

2 Data and Input Access

Item	Current evidence
Source and working link	Self-authored level prompts and passwords plus logged attempts (https://github.com/rickeyjohnson/Locksmith); model qwen3:8b served locally by Ollama; 60 of the 78 attack prompts sampled from the public Tensor Trust dataset (Toyer et al., ICLR 2024) (https://github.com/rickeyjohnson/Locksmith/blob/main/docs/research/literature.md).
Access and amount	All 8 levels run, and 1,170 benchmark attempts are logged: 624 against the stacked ladder and 546 against single defenses.
Format and preparation	One labeled row per attempt in Supabase Postgres (local JSONL fallback); leaked / not leaked is assigned automatically by a detector covering spaced, reversed, base64, hex, and acrostic spellings, and failed model calls are excluded.
Concerns	Players consent before playing and are told not to enter personal information; passwords are generated nonsense words kept out of the public repository, and row-level security blocks all public database access.

3 Working Technical Artifact

Artifact evidence	Response
Artifact type	Web application (Next.js + TypeScript) plus a command-line attack benchmark.
File or working link	https://github.com/rickeyjohnson/Locksmith — runs locally per the README (‘npm run dev’); public deployment before the final submission.
What currently works	A typed prompt runs through input guards, the level's system prompt and the model, then output guards, producing a guardian reply, an automatic leak label, a logged row, and level unlocking on a correct guess; all 8 levels work and 87 tests pass.
Current limitation	No human playtest data yet, so attempts-until-success is unmeasured, and levels 7-8 may be unbeatable since no attack in the corpus leaked past them.



4 Baseline or Initial Result

Method	Input or sample	Initial output	What it demonstrates
Stacked ladder, levels 1-8, against the non-AI keyword and output filters	78 attack prompts per level; qwen3:8b, temperature 0.7, seed 42 (624 attempts)	Player-visible leak rate: 84.6%, 47.4%, 6.4%, 0.0%, 1.4%, 0.0%, 0.0%, 0.0%; model-side rate once the prompt reaches the model: 84.6%, 82.2%, 86.7%, 60.0%, 21.6%, 2.2%, 0.0%, 0.0%.	Leak rate falls as complexity rises, but the keyword filter blocks 42% of traffic without changing the model's behavior at all (86.7% vs 84.6% undefended), so only prompt-level defenses make the model safer; the measure saturates at 0% from level 6 up.

5 Evaluation Readiness

Identify at least two measures from the approved proposal and show how each will be used.

Measure	What will be evaluated	Comparison or target	Readiness
Measure 1: password-leak rate per level	Share of attempts whose shown response contains the password in any disguise	AI levels vs the non-AI filter baseline; target is a rate falling level to level	Ready — produced for all 8 levels and 7 isolated defenses
Measure 2: mean and median attempts until success	Attempts a player makes before guessing the password correctly	Compared across levels; more attempts means a stronger defense	In progress — logging built; awaiting the Sep 26-27 playtest

6 Progress Status

Project component	Status	Evidence
Data or input access	Complete	8 level configs, a 78-prompt attack corpus, and qwen3:8b verified locally
Data preparation	Complete	Automatic leak labeling (11 unit tests) over 1,170 rows in docs/results/
Baseline	Complete	Non-AI baselines measured alone: keyword filter 50.0%, output filter 3.8%, against 84.6% undefended
AI or ML component	Complete	All 8 levels call one pinned model; AI prompt shield and critic agent implemented and tested
Evaluation	In progress	Leak rate automated; attempts-until-success awaits playtesters (scripts/analyze.ts)
Final artifact	In progress	Playable 8-level game with consent, logging, and 87 tests; deployment remains

7 Risks and Next Steps

Prompt	Response
Most important completed task	A tested game and benchmark measuring seven defenses on one fixed model, showing that input filtering reduces exposure without making the model safer.
Greatest current blocker or risk	The measure saturates: levels 6-8 sit at 0%, so the AI defenses cannot be ranked and playtesters may stall before reaching them.
Before Practical Checkpoint 2	Run the Sep 26-27 playtest, report human leak rates and attempts-until-success per level, and write the final evaluation report.
Backup plan	Report the benchmark as primary evidence if playtesting falls short, and collapse to five levels if the upper ladder proves unusable.

Submission Check

- ☒ ☐ The checkpoint report is no more than two pages and is exported as a PDF.
- ☒ ☐ The current technical artifact is submitted separately or linked with working access.
- ☒ ☐ The artifact accepts a real sample input and produces an observable output.
- ☒ ☐ The report includes at least one baseline or initial result.
- ☒ ☐ All links open without requesting additional permission.
- ☒ ☐ No private, protected, confidential, or personally identifiable data is included.

Student name **Rickey Johnson**

Date 09/20/2026